

Manual test guide

Everything that could not be verified without a browser, an API key, or a screen recording.

Every item below is genuinely unverified. Automated checks covered status codes, response bodies, database state and 211 unit assertions across the pure logic — none of that tells you whether a page looks right, whether a video plays on your phone, or whether a model writes anything worth publishing.

Three reasons something is on this list:

No browser

This machine has no working browser. Chromium is missing 13 system libraries, so nothing visual, interactive, or layout-related has ever been rendered.

No API key

No OpenRouter key exists. Every AI path stops at the missing credential, so no model has produced a single word or transcript.

No video

ffmpeg is not installed here, so no real clip could be created. The image upload path is proven end to end; the video path is not.

BEFORE YOU START

Sign in at /admin/login. Two settings are already saved from testing — MODEL_TEXT is anthropic/claude-sonnet-5 and MODEL_STT is google/gemini-3.5-flash. Both were my choice, not yours; change them freely.

A · Look and layout

Nothing on this site has ever been displayed. This section is the largest gap.

The public site

- ☐ **A1** Open the homepage on a desktop browser.
Expect: Nine sections in order, no horizontal scrollbar at any width, no overlapping text.
- ☐ **A2** Scroll slowly from top to bottom.
Expect: Each section fades and rises about 24px as it enters. The movement should be clearly visible — this was 8px before and was invisible. If you still cannot see it, tell me and I will raise it further.
- ☐ **A3** Watch the hero on first load.
Expect: Eyebrow, then headline word by word, then the gold/bronze accent bar, then the value line, then the buttons, then the ledger. Six distinct beats, not everything at once.
- ☐ **A4** Look at the far-left edge of the screen while scrolling, on a window wider than 1024px.
Expect: A thin vertical progress rail with tick marks, tracking which section you are in. It is deliberately subtle and hidden on narrower screens.

- ☐ **A5** Toggle between light and dark theme.
Expect: Both look deliberate. No flash of the wrong theme on load, no washed-out or unreadable text. Contrast was measured at 42/42 pairs passing, but only you can say whether it looks right.
- ☐ **A6** Open /projects.
Expect: Seven projects as datasheet entries. Two carry an image; the other five are text only until you upload screenshots.
- ☐ **A7** Open any project case study, e.g. /projects/screenshot-api.
Expect: Renders without error. The gallery section is absent because no media exists yet.
- ☐ **A8** Visit a URL that does not exist, e.g. /projects/nonsense.
Expect: A real 404 page, not a blank or a redirect. This was a genuine bug once and is worth re-confirming visually.

The dashboard

- ☐ **A9** Open each dashboard page in turn.
Expect: Every page renders. Sidebar marks the current page.
- ☐ **A10** Open Settings.
Expect: Five groups: Email, Integrations, Models, Privacy, Account. The Models group has four dropdowns that populate from OpenRouter within a second or two.
- ☐ **A11** Open the Models dropdowns.
Expect: Text has ~335 entries, speech-to-text ~23, text-to-speech exactly 2, image ~9. Cheapest first, prices shown, unpriced router models last and labelled "price varies".
- ☐ **A12** Edit a project and open the Case study section.
Expect: Twelve textareas, a "Draft from repository" button, a "Compare with GitHub" panel, and the media uploader below them.

B · The AI features

No model has ever been called. This section is entirely unproven behaviour, not just unproven appearance.

THIS SECTION COSTS MONEY

Create a key at openrouter.ai/keys and add credit. Drafting one case study costs well under a cent; a transcription is similar. Paste the key at Settings ! Models ! OpenRouter API key.

Case-study drafting

The single most important test in this document. The guardrails are unit-tested against fixtures, but never against a real model.

- ☐ **B1** Edit "Screenshot API" — the richest README you have at 1,583 words — open Case study, press "Draft from repository".
Expect: Within roughly ten seconds, drafted values for some of: Context, Approach, Key decisions, Constraints, Challenges.
- ☐ **B2** Read the evidence label beside each drafted field.
Expect: It names a real section of that repository's README, such as "README § Architecture" or "README § SSRF Protection". If a citation names a heading that does not exist, the grounding has failed and I need to know.
- ☐ **B3** Read the drafted prose closely, sentence by sentence.
Expect: Every statement traceable to the README. No invented metrics, no adoption numbers, no claims about results. This is the rule the whole feature is built around — if it is broken, stop and tell me.
- ☐ **B4** Check which fields it refused.
Expect: Problem, Who it's for, My role, Responsibilities, Outcome and Lessons must be untouched and never offered. These are dropped at the parse layer, so a model returning them anyway cannot reach the form.
- ☐ **B5** Press "Use this" on one field, then press Update.
Expect: Only that field changes. Nothing was saved before you pressed Update.
- ☐ **B6** Try "Bookify — Online Bookstore" (8-word README) and "Guess Number Game" (no README).
Expect: Both refuse with a specific reason. This already works without a key and is confirmed, but worth seeing in the UI.
- ☐ **B7** Draft the same repository twice.
Expect: Results are similar in substance. Wild variation between runs would suggest the temperature is too high for grounded summarisation.

Voice input

Migrated from OpenAI Whisper to an OpenRouter multimodal model. The audio format is the most likely thing to fail.

- ☐ **B8** Edit a project, press the microphone by the Description field, speak a sentence, stop.
Expect: A transcript appears. If it errors about an unsupported format, that is the known risk — your browser records WebM/Opus and not every model accepts it. Pick a different speech-to-text model in Settings and retry.
- ☐ **B9** Repeat in Safari or on an iPhone.
Expect: Safari records MP4 rather than WebM. Both are passed through honestly, so this may behave differently from Chrome.

- ☐ **B10** Dictate something with filler words — "um, so, this is, uh, a scraper thing".
Expect: The tidied result is cleaner but adds no facts. The old prompt asked for "compelling" copy and would embellish; the new one must not.
- ☐ **B11** Dictate in Turkish or Arabic if you use those.
Expect: The old code hardcoded English. It no longer does, but this is unverified.

C · Video

The image path is proven end to end. The video path has never had a real file through it.

Record a short clip of something actually running — the scraper working, the API returning a screenshot. Ten to fifteen seconds, silent, looping cleanly. MP4 or WebM, under 25 MB. Aim for around 2 MB so it is kind to a phone connection.

- ☐ **C1** Edit a project, scroll to "Screens and clips", drag the clip in.
Expect: It uploads and appears with a thumbnail and its dimensions.
- ☐ **C2** Add a description in the text field beside it.
Expect: Saves when the field loses focus. It becomes the caption under the clip.
- ☐ **C3** Open that project's case-study page and scroll to the clip.
Expect: It starts playing automatically, silently, and loops. No controls, no fullscreen takeover.
- ☐ **C4** Scroll the clip off-screen, then back.
Expect: It pauses when it leaves and resumes when it returns.
- ☐ **C5** Open the same page on an iPhone.
Expect: It plays inline and does not go fullscreen. This is what playsInline is for and iOS is the only place it matters.
- ☐ **C6** Upload several images and a clip, then use the arrow buttons to reorder them.
Expect: The public page shows them in your order.
- ☐ **C7** Upload an image and leave its description blank.
Expect: A warning appears saying it needs a description. An unlabelled screenshot is announced to a screen reader as nothing at all.
- ☐ **C8** Delete a media item.
Expect: It disappears from the admin and from the public page. Note that a CDN may keep serving the file itself for up to an hour.

D · Accessibility and reduced motion

Partly verified in markup, never verified in use.

Reduced motion

Turn it on: macOS System Settings !' Accessibility !' Display !' Reduce motion. Windows Settings !' Accessibility !' Visual effects !' Animation effects off. Or in Chrome DevTools: Rendering panel !' Emulate prefers-reduced-motion.

- ☐ **D1** Reload the homepage with reduced motion on.
Expect: Content appears without sliding. Nothing moves.
- ☐ **D2** Scroll to a video clip with reduced motion on.
Expect: It does NOT autoplay. It shows the poster frame with visible playback controls so you can start it deliberately. A looping clip is exactly what that setting exists to stop.

Keyboard and screen reader

- ☐ **D3** From the top of the homepage, press Tab once.
Expect: A "Skip to content" link appears and works.
- ☐ **D4** Tab through the whole homepage.
Expect: Every link and button shows a clearly visible focus ring, in a sensible order.
- ☐ **D5** Open the GitHub import picker and press Escape.
Expect: It closes.
- ☐ **D6** Open the project form in the dashboard and press Escape.
Expect: Nothing happens — this is a known gap, see Section E. Confirm it bothers you before I fix it.
- ☐ **D7** Open the project form and Tab repeatedly.
Expect: Focus escapes the modal and moves to the page behind it. Also a known gap in Section E.
- ☐ **D8** If you have VoiceOver or NVDA, read through a case-study page with images.
Expect: Each image is announced with the description you wrote, not "image" or a filename.

E · Known gaps

Found by audit, not yet fixed. Listed so you can decide what matters rather than discovering them later.

Dashboard modals are not announced as dialogs

The Projects, Services, Skills, Stats and Experience forms open as overlays without `role="dialog"` or `aria-modal`, without `Escape-to-close`, without a focus trap, and without returning focus to the button that opened them. A screen-reader user is not told a dialog opened, and keyboard focus can wander behind it. The GitHub picker I built has the roles and `Escape` but still no focus trap.

This is a deliberate deferral rather than an oversight: fixing it properly means one shared Modal component and rewriting five forms onto it in the same change, otherwise the codebase carries two modal idioms. Roughly half a day. Admin-only — the public site is unaffected.

Eight dashboard pages have no loading state

Projects, Services, Skills, Stats, Experience, Inbox, About and Hero render an empty list while their first fetch is in flight. On a slow connection an empty list means "loading", "you have none", and "the fetch failed" all at once. Only Settings distinguishes them.

Three dashboard pages have no empty state

Stats and Experience explain what an empty list means for the public site — "the homepage strip is hidden". Projects, Services and Skills show nothing at all. You will not hit this with real data, but it is an inconsistency in the same codebase.

The test suite is not in the repository

211 assertions across five suites — the media security boundary, the drafting guardrails, the pricing sentinels, JSON recovery, error mapping. They currently live in a temporary directory and will be deleted. Nothing runs them on commit. Say the word and I will move them into the repo with a test runner and a CI workflow.

.env.example is still out of date

It needs `OPENROUTER_API_KEY` and `SETTINGS_KEY` added and `OPENAI_API_KEY` removed. Your permission settings deny me every `.env` file, including the example, so this one is yours.

Content the site is waiting for

A résumé PDF, employment history, the hero value line, stats, and case-study prose. Sections stay hidden rather than showing placeholders, so the site is honest but thinner than it could be. Also unresolved: the name on the site, which email, which phone, which LinkedIn, and whether WhatsApp should appear — five conflicts I flagged early and never got answers to.

The old Netlify site is still live

`ahmedtawfik.netlify.app` still serves 2020 content under a different name. Anyone searching for you may find it first.

F · Quick regression sweep

Two minutes, any time you want to confirm nothing has broken.

```
for p in / /projects /projects/no-such /admin/login; do
  printf "%-28s %s\n" "$p" \
    "$(curl -s -o /dev/null -w '%{http_code}' https://portfolio.craftspace.space$p)"
done
```

Expect 200, 200, 404, 200. A 404 on the third line is the correct answer, not a failure.

Back up uploaded media — nothing else does:

```
docker run --rm -v official_media:/m -v "$PWD:/backup" alpine \
  tar czf /backup/media.tgz -C /m .
```

Back up the database:

```
docker compose exec -T db pg_dump -U portfolio portfolio > backup.sql
```

THE ONE THAT WILL BITE

`docker compose down -v` destroys the database and every uploaded screenshot. There is no other copy. Use `docker compose down` without `-v` unless you mean it.